

# CS549 Distributed Systems & Cloud Computing

Dominic Duggan  
Stevens Institute of Technology

1

**THIS COURSE**

2

## Graduate Certificate / MS in Enterprise and Cloud Computing

- CS522 Mobile Systems and Apps
- CS526 Enterprise and Cloud Computing
- CS548 Enterprise Software Architecture
- **CS549 Distributed Systems and Cloud Computing**
- CS594 Enterprise and Cloud Security

3

## MS/ECC Program Outcomes

	CS522	CS526	CS548	CS549	CS594
Infrastructure		X Virtualization		X Availability	
Data Modeling			X		
Design			X		
Applications	X	X	X	X	
Security & Privacy		X Secure Virt			X

4

## Web Apps & Web Services

	Azure (CS526)	Java (CS548, CS549)
Client-Side B2C	AJAX Silverlight	AJAX Android
Server-Side B2C	ASP.NET	Java Faces
Server-Side B2B	Windows Communication Foundation (WCF)	<b>Java EE (CDI, JCA)</b> <b>Jersey</b> <b>Atmosphere</b>
Database	Entity Framework LINQ	Java Persistence Architecture (JPA)

## Cloud Computing

Dominic Duggan

Stevens Institute of Technology

Based on material by K. Birman, P. Francis, A.M.  
Nguyen, A. Tanenbaum

# What is Cloud Computing?

- *I don't understand what we would do differently in the light of Cloud Computing other than change the wordings of some of our ads*

Larry Ellison, Oracle's CEO

- *I have not heard two people say the same thing about it [cloud]. There are multiple definitions out there of "the cloud"*

Andy Isherwood, HP's Vice President of European Software Sales

- *It's stupidity. It's worse than stupidity: it's a marketing hype campaign.*

Richard Stallman, Free Software Foundation founder

7

## SOFTWARE AS A SERVICE AND UTILITY COMPUTING

8

## Software as a Service (SaaS)

Ex: Salesforce.com

### Traditional Software



**Build Your Own**

### On-Demand Utility



**Plug In, Subscribe  
Pay-per-Use**

9


## Software as a Service (SaaS)



- **Application used as on demand service**
  - Often provided via the Internet
- Example: Google Apps
- Benefits to users
  - Reduce expenses: multiple computers, multiple users
  - Ease of usage: easy installation, access everywhere
- Benefits to providers
  - Easier to maintain
  - Control usage (no illegal copies)

10

# Utility Computing (UC)

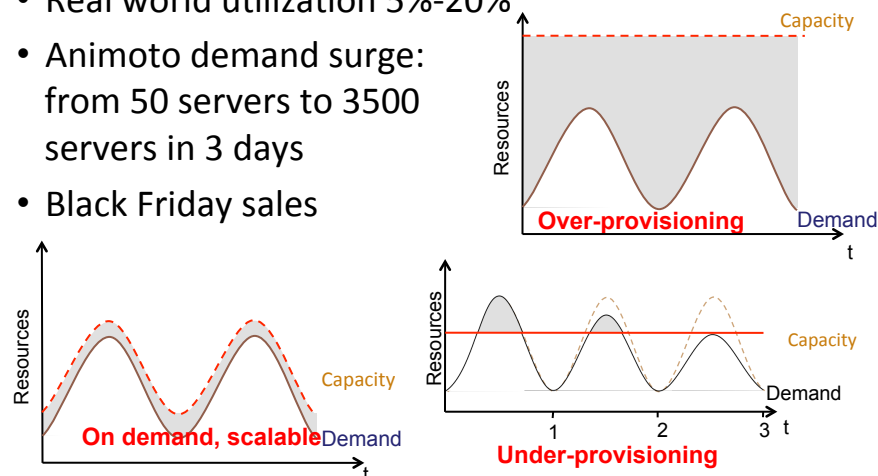


- **Computing resources** (cpu hours, memory, network) and platform to run software are **provided as on demand service**
  - *Hardware as a service (HaaS)*
  - *Infrastructure as a service (IaaS)*
  - *Platform as a Service (PaaS)*
- Examples of UC providers:
  - PaaS: MS Azure ...
  - IaaS: Amazon EC2 ...
- Who will use UC?

11

## Utility Computing: Mitigate Risks

- Real world utilization 5%-20%
- Animoto demand surge: from 50 servers to 3500 servers in 3 days
- Black Friday sales



12

## Utility Computing – Amazon EC2

- **Elastic Compute Cloud**
- Rent VM instances to run your software
- Full root-level access to VM

13

## Amazon EC2

- Create an Amazon Machine Image (AMI)
- Upload AMI to Amazon S3 (simple storage service)
- Use Amazon EC2 web service to configure
- Choose OS, start AMI instances

PHP  
Apache  
Perl  
Postgress  
Linux-Ubuntu

Ruby  
Rails  
MySQL  
Fedora-6

WebSphere  
Hibernate  
Java  
Linux


**Amazon S3**

**Amazon EC2**

14

## Amazon EC2


- Create an Amazon Machine Image (AMI)
- Upload AMI to Amazon S3 (simple storage service)
- Use Amazon EC2 web service to configure
- Choose OS, start AMI instances



15

## Amazon EC2

- Create an Amazon Machine Image (AMI)
- Upload AMI to Amazon S3 (simple storage service)
- Use Amazon EC2 web service to configure
- Choose OS, start AMI instances



16




## Utility Computing – MS Azure

- Write your web program and submit to Azure
- How to use
  - Download MS SDK, Azure tools
  - Develop your program locally
  - Register for an Azure id
  - Launch your application in Azure

17

## Spectrum Of Abstractions


- Different levels of abstraction
  - Instruction Set VM: Amazon EC2
  - Framework VM: MS Azure
- Similar to languages
  - Higher level abstractions can be built on top of lower ones



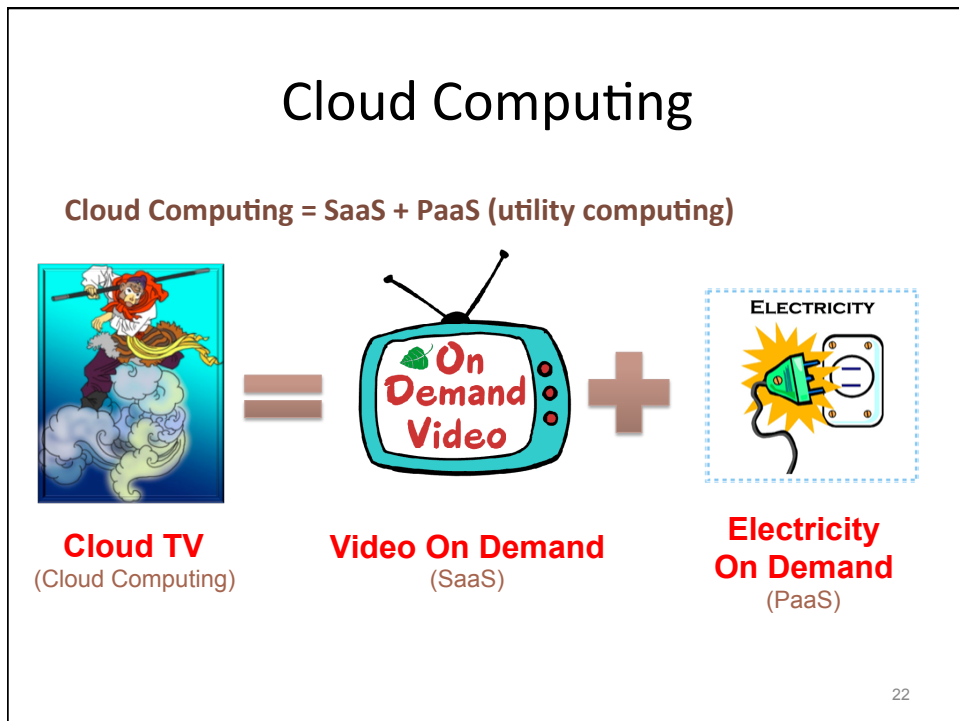
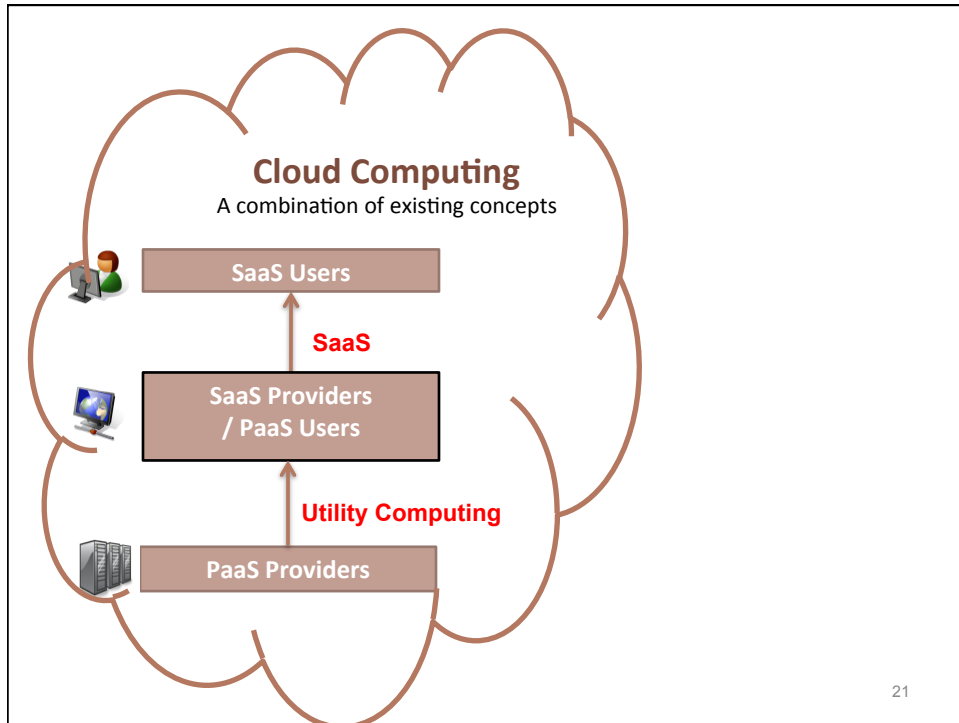
18

# CLOUD COMPUTING

19



20



## Significance of Cloud Computing

- The illusion of infinite computing resources
- The elimination of an up-front commitment by users
- The ability to use and pay on demand
- Cloud Computing vs P2P?
  - Both take advantage of remote resources
  - P2P: does not use clouds (datacenters), peers do not get paid, lower reliability
- Cloud Computing vs Grid Computing?
  - Both use clouds
  - Grid Computing requires commitment, share based on common interests. Not public cloud

23

## Cloud Killer Apps

- Mobile and web applications
  - Mobile devices: low memory & computation power
- Extensions of desktop software
  - Matlab, Mathematica
- Batch processing / MapReduce
  - Peter Harkins at The Washington Post: 200 EC2 instances (1,407 server hours), convert 17,481 pages of Hillary Clinton's travel documents within 9 hours
  - The New York Times used 100 Amazon EC2 instances + Hadoop application to recognize 4TB of raw TIFF image into 1.1 million PDFs in 24 hours (\$240)

24

## ECONOMICS OF THE CLOUD

25

### Should I Move Into A Cloud?

- Does it really save money?

$$\text{UserHours}_{\text{cloud}} \times (\text{revenue} - \text{Cost}_{\text{cloud}}) \geq \text{UserHours}_{\text{datacenter}} \times \left( \text{revenue} - \frac{\text{Cost}_{\text{datacenter}}}{\text{Utilization}} \right)$$

- $\text{Cost}_{\text{cloud}} > \text{Cost}_{\text{datacenter}}$ , balance by *Utilization*
- $\text{UserHours}_{\text{cloud}} > \text{UserHours}_{\text{datacenter}}$  (under-provisioning)

- Other factors
  - Re-implement programs
  - Move data into cloud
  - What else?
- Example:
  - Upload rate 20Mbits / s. 500GB takes 55 hours
  - If can process locally in less than 55 hours → moving into a cloud would not save time

26


## Adoption Challenges

Challenge	Opportunity
Availability	Multiple providers
Data lock-in	Standardization
Data Confidentiality and Auditability	Encryption, VLANs, Firewalls

- Coghead, a cloud vendor closed its business in February 2009
  - Customers need to rewrite their applications
  - Another company will automatically convert customer data to their proprietary formats...
- Online storage service The Linkup closed July 10, 2008
  - 20,000 paying subscribers lost their data

27

## Adoption Challenges



28

Cloud Control, InformationWeek Reports, 2009


## Growth Challenges

Challenge	Opportunity
Data transfer bottlenecks	FedEx-ing disks, reuse data multiple times
Performance unpredictability	Improved VM support, flash memory
Scalable storage	Invent scalable storage
Bugs in large distributed systems	Invent Debugger using Distributed VMs
Scaling quickly	Invent Auto-Scaler

29

## Growth Challenges

- Data transfer bottle neck
  - WAN cost reduces slowest:  
2003 → 2008: WAN 2.7x, CPU 16x, storage 10x
  - Fastest way to transfer large data: send the disks
- Performance unpredictability
  - Large variation in I/O operations
  - Inefficiency in I/O virtualization



## Policy And Business Challenge

Challenge	Opportunity
Reputation Fate Sharing	Offer reputation-guarding services like those for email
Software Licensing	Pay-for-use licenses; Bulk use sales

- Reputation: Many blacklists use IP addresses and IP ranges
- Software licensing:
  - Open source software readily applicable
  - Windows, IBM software offered per hour for EC2

31

## ECONOMICS OF THE CLOUD

32



## Should I Move Into A Cloud?

- Does it really save money?

$$\text{UserHours}_{\text{cloud}} \times (\text{revenue} - \text{Cost}_{\text{cloud}}) \geq \text{UserHours}_{\text{datacenter}} \times \left( \text{revenue} - \frac{\text{Cost}_{\text{datacenter}}}{\text{Utilization}} \right)$$

- $\text{Cost}_{\text{cloud}} > \text{Cost}_{\text{datacenter}}$ , balance by *Utilization*
- $\text{UserHours}_{\text{cloud}} > \text{UserHours}_{\text{datacenter}}$  (under-provisioning)

- Other factors
  - Re-implement programs
  - Move data into cloud
  - What else?
- Example:
  - Upload rate 20Mbits / s. 500GB takes 55 hours
  - If can process locally in less than 55 hours → moving into a cloud would not save time

33


## Adoption Challenges

Challenge	Opportunity
Availability	Multiple providers
Data lock-in	Standardization
Data Confidentiality and Auditability	Encryption, VLANs, Firewalls

- Coghead, a cloud vendor closed its business in February 2009
  - Customers need to rewrite their applications
  - Another company will automatically convert customer data to their proprietary formats...
- Online storage service The Linkup closed July 10, 2008
  - 20,000 paying subscribers lost their data

34

## Adoption Challenges



Cloud Control, InformationWeek Reports, 2009

35


## Growth Challenges

Challenge	Opportunity
Data transfer bottlenecks	FedEx-ing disks, reuse data multiple times
Performance unpredictability	Improved VM support, flash memory
Scalable storage	Invent scalable storage
Bugs in large distributed systems	Invent Debugger using Distributed VMs
Scaling quickly	Invent Auto-Scaler

36

## Growth Challenges

- Data transfer bottle neck
  - WAN cost reduces slowest:  
2003 → 2008: WAN 2.7x, CPU 16x, storage 10x
  - Fastest way to transfer large data: send the disks
- Performance unpredictability
  - Large variation in I/O operations
  - Inefficiency in I/O virtualization



## Policy And Business Challenge

Challenge	Opportunity
Reputation Fate Sharing	Offer reputation-guarding services like those for email
Software Licensing	Pay-for-use licenses; Bulk use sales

- Reputation: Many blacklists use IP addresses and IP ranges
- Software licensing:
  - Open source software readily applicable
  - Windows, IBM software offered per hour for EC2

# DISTRIBUTED SYSTEMS

39

## Some terminology

- A **program** is the code you type in
- A **process** is what you get when you run it
- A **message** is used to communicate between processes.
- A **packet** is a fragment of a message that might travel on the wire.
  - Variable size
  - Limited size
- A **protocol** is an algorithm by which processes cooperate to do something using message exchanges.


40

## More terminology

- A **network** is the infrastructure that links the computers, etc.
  - routers
  - communication links
- **Network application**: fetches needed data from servers over the network
- **Distributed system**: multiple processes that cooperate to do something

41

A network is like a “mostly reliable” post office



42

## Loss of reliability

- Links can corrupt messages
  - Internet “backbone”
  - Wireless connections, cable modems, ADSL
- Routers can get overloaded
- Solution: retransmission protocols

43

## Distributed systems vs network applications

- **Distributed systems**
  - many components
  - often mimic a single, non-distributed process
- **Networked application**
  - centered around the user or computer where it runs


44

## What about the Web?

- Browser is independent
- REST: Web servers don't keep track of clients.
  - Cookies
  - Database of account info
- Two network applications that talk to each other


45

## What about the Web?



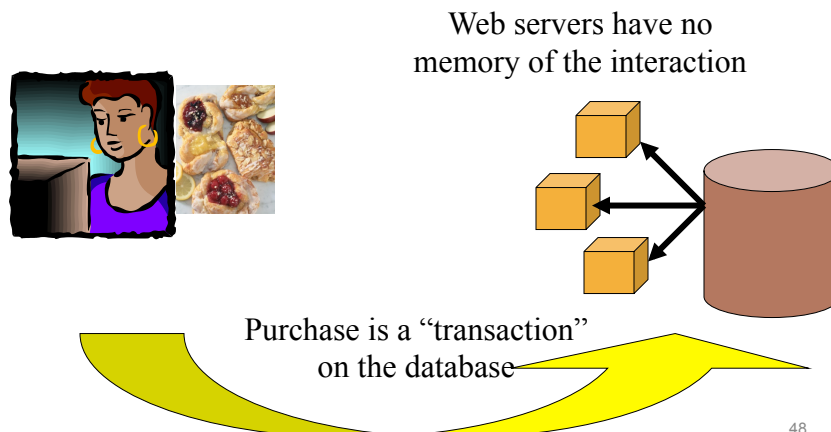
46

## What about the Web?



47

## What about the Web?



48



## What about the Cloud?

- Data center or cloud is a complex distributed system
  - Many servers
  - Routing clients to servers
  - Data replicated
    - load balancing
    - high availability
  - Complex security and administration policies
- **“Network application”** talking to a **“distributed system”**

49

**NETWORKS**

50

## Who recognizes this?

```
int sockfd;  
struct sockaddr_in addr;  
  
addr.sin_family = AF_INET;  
addr.sin_addr.s_addr =  
    inet_addr(SERV_HOST_ADDR);  
addr.sin_port = htons(SERV_TCP_PORT);  
  
sockfd = socket(AF_INET, SOCK_STREAM, 0);  
connect(sockfd, (struct sockaddr *) &addr,  
        sizeof(serv_addr));  
do_stuff(stdin, sockfd);
```

51

## Classic view of network API

- Start with host name  
(maybe)

foo.bar.com

52

## Classic view of network API


- Start with host name
- Get an IP address



53

## Classic view of network API


- Start with host name
- Get an IP address
- Make a socket (protocol, address)



54

## Classic view of network API

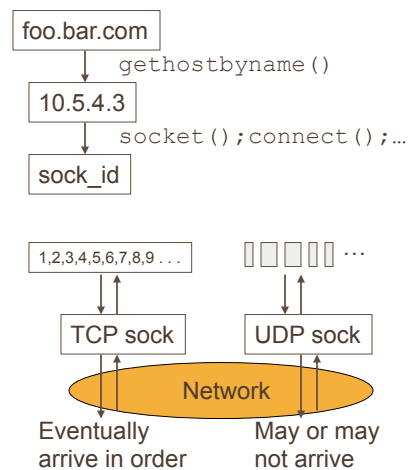
- Start with host name
- Get an IP address
- Make a socket (protocol, address)
- Send byte stream (TCP)



55

## Classic view of network API

- Start with host name
- Get an IP address
- Make a socket (protocol, address)
- Send byte stream (TCP) or packets (UDP)



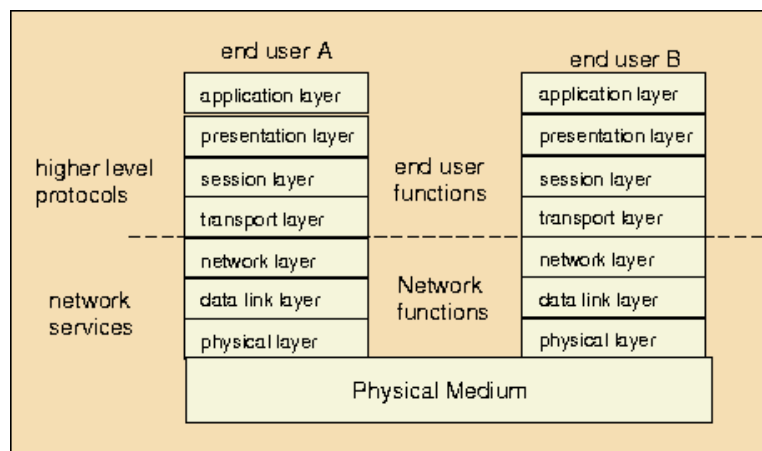
56

## Classic approach “broken” in many ways

- **DNS:** IP address different depending on who asks
- **NAT:** Address may be changed in transit
- **Firewall:** IP address may not be reachable
  - Or may be reachable by you but not another host
- **DHCP:** IP address may change
- **Caches:** Packets may not come from who you think

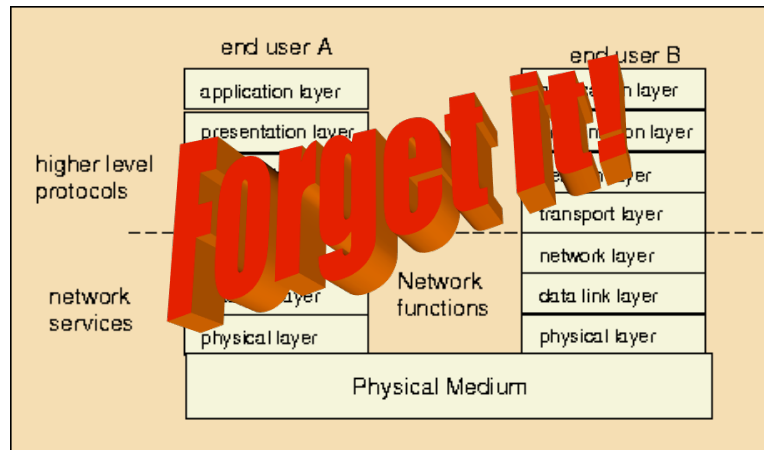
57

## Classic OSI stack



58

## Classic OSI stack




59

## TUNNELING

60

# Tunneling Analogy

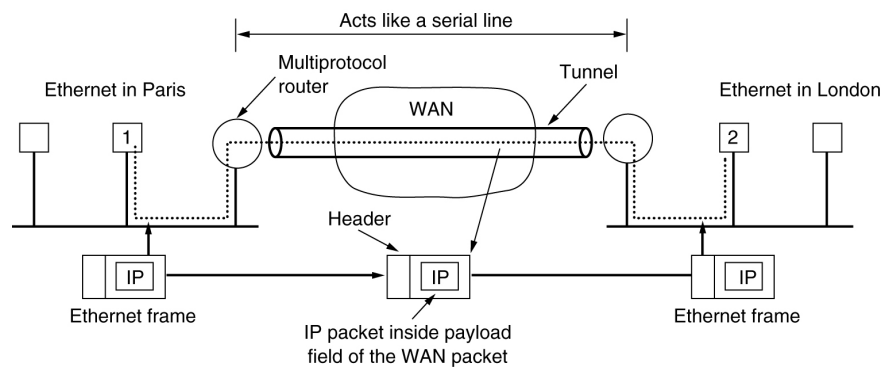
Tunneling a car from France to England.



61

# Tunneling in Networks

Tunneling a packet from Paris to London.



62

## Example Microsoft VPN stack

Application
TCP
IP
PPP
L2TP
UDP
IPsec
IP
PPP
PPPoE
Ethernet

63

## Example Microsoft VPN stack

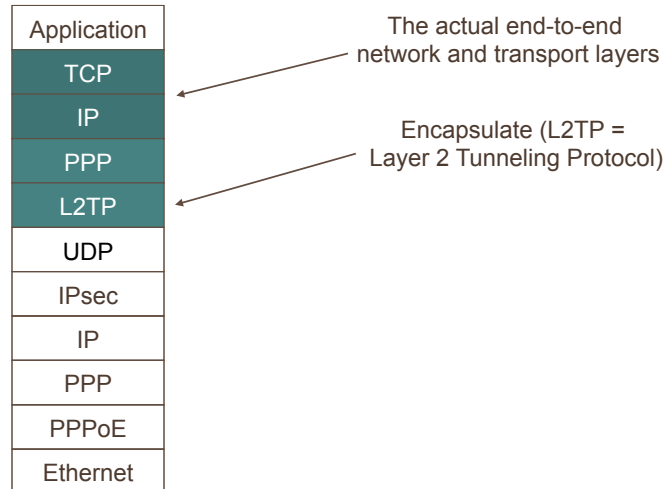
Application
TCP
IP
PPP
L2TP
UDP
IPsec
IP
PPP
PPPoE
Ethernet

The actual end-to-end  
network and transport layers

64

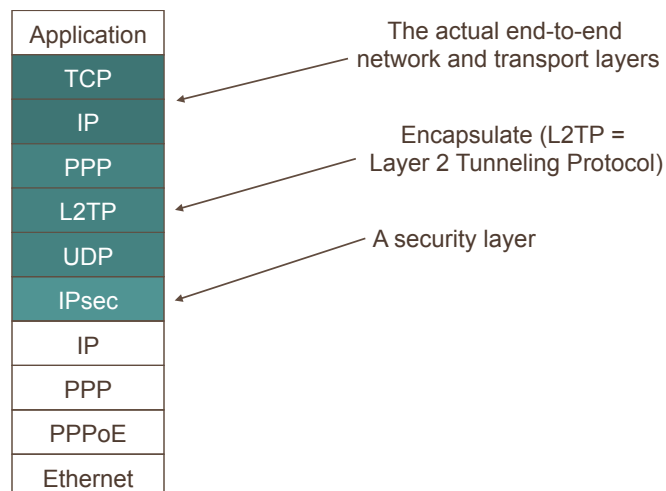


## Example Microsoft VPN stack



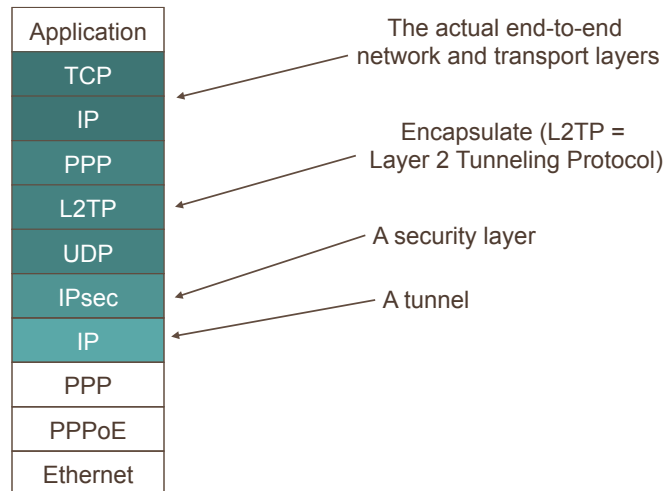
65

## Example Microsoft VPN stack



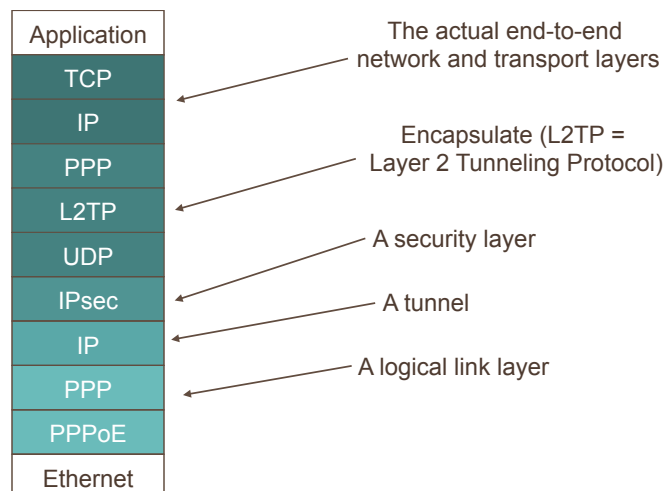
66

## Example Microsoft VPN stack



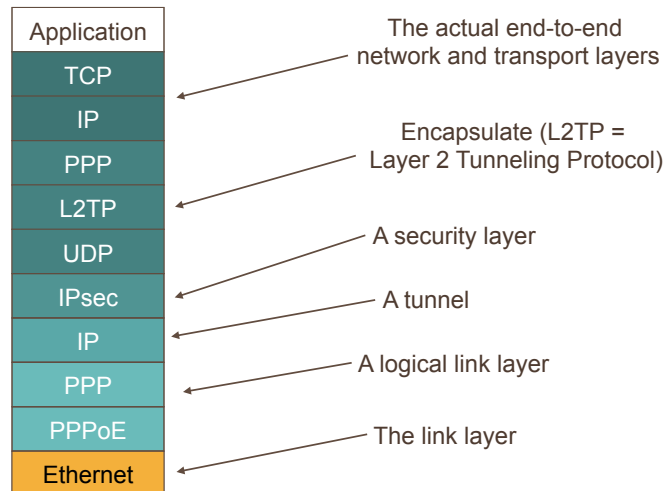
67

## Example Microsoft VPN stack



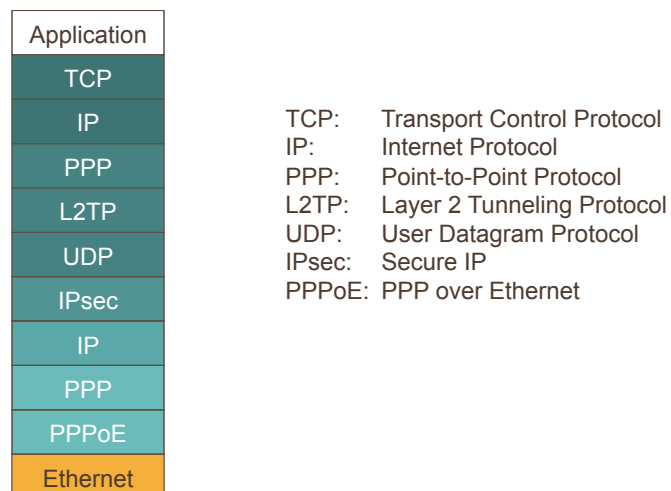
68

## Example Microsoft VPN stack



69

## Example Microsoft VPN stack



70

## END-TO-END ARGUMENT


71

## End-to-End argument

- Internet:
  - Suppose an IP packet will take  $n$  hops
  - probability  $p$  of loss on each hop
  - Transfer a file of  $k$  IP packets
  - Should we:
    - use a retransmission protocol running “end-to-end” or
    - $n$  TCP protocols in a chain?

72

## End-to-End argument



Probability of successful transit:  $(1-p)^n$ ,  
Expected packets lost:  $k - k \cdot (1-p)^n$

73

## Saltzer et. al. analysis

- If  $p$  is very small, then even with many hops most packets will get through
  - Overhead of using TCP protocols in the links
  - End-to-end recovery mechanism

74

## Generalized End-to-End view?

- Low-level mechanisms should focus on speed, not reliability
- The application should worry about “properties” it needs
- In general, add additional functionality *end-to-end in the application*
  - Not in the network

75

## What can we learn from this?



- That the internet is a mature technology
  - Kludges on kludges
- *That the end-to-end argument actually works!*

76

## When should the network do more?


- When you get performance gains
  - Link-level retransmissions over a lossy link
  - Ex: wireless network
- Also
  - When the network doesn't trust the end user
    - Corporation or military
  - Some things can't be done at the end
    - Routing algorithms
    - Billing
    - User authentication

77

## NETWORK INFRASTRUCTURE

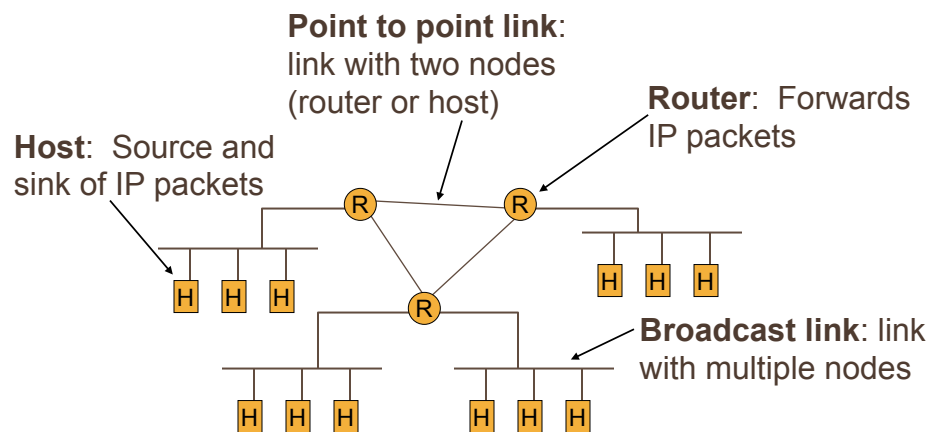
78

A network is like a “mostly reliable” post office



79

## Network components



80




# Network components

- **Network:** Collection of hosts, links, and routers
- **Site:** Stub network, typically in one location and under control of one administration
- **Firewall/NAT:** Box between the site and ISP
- **ISP:** Internet Service Provider. Transit network that provides IP connectivity for sites
- **Backbone ISP:** Transit network for regional ISPs and large sites
- **Inter-exchange (peering point):** Broadcast link where multiple ISPs connect and exchange routing information (peering)
- **Hosting center:** Stub network that supports lots of hosts (web services), high speed connections to many backbone ISPs.
- **Bilateral peering:** Direct connection between two backbone ISPs

81

# Internet topology



82

# NETWORK PROTOCOLS


83

## Protocol layering

- Communications stack consists of a set of **services**
  - Each providing a service to the layer above
  - Using services of the layer below
  - Each service has a programming API
- Each service has to convey information for one or more **peers** across the network
- This information is contained in a **header**

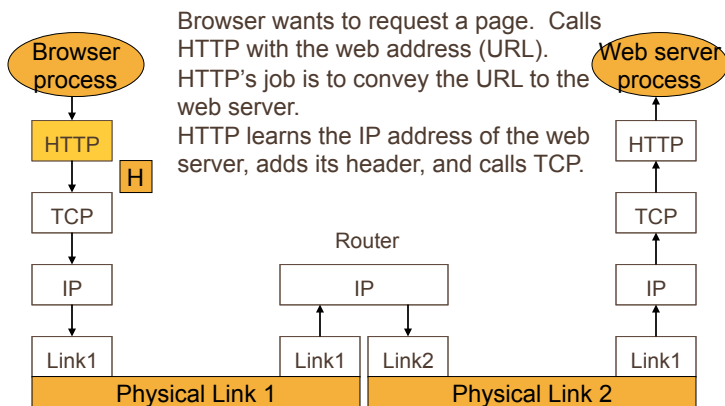
84

## Protocol layering example



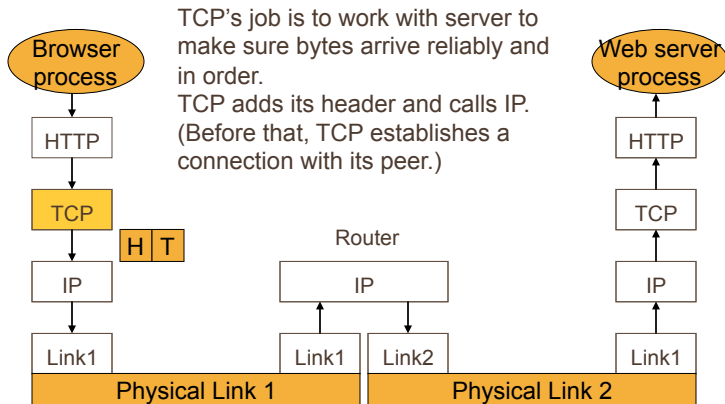
85

## Protocol layering example



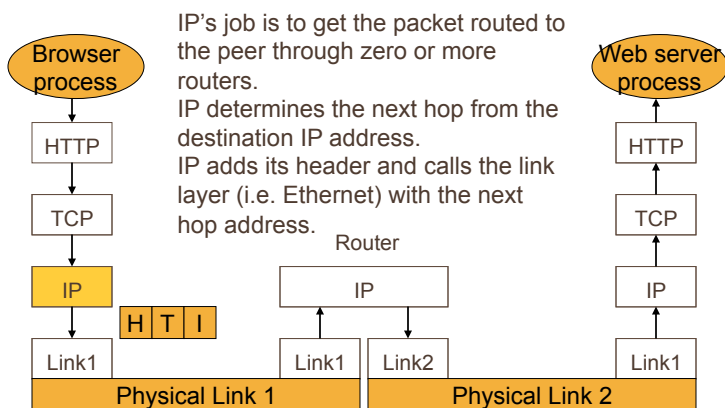
86

## Protocol layering example



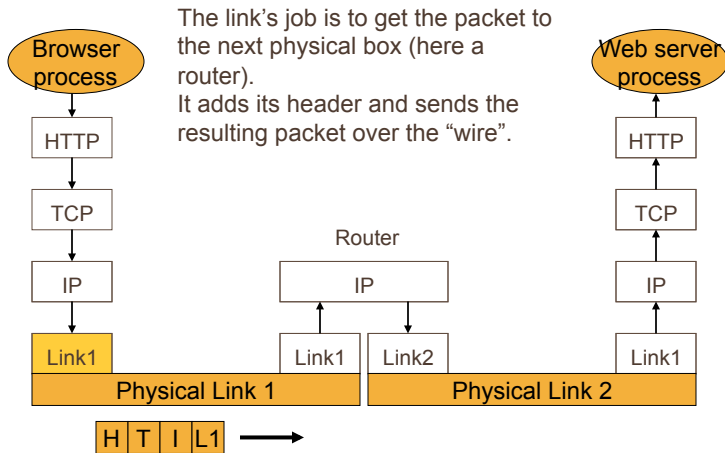
87

## Protocol layering example



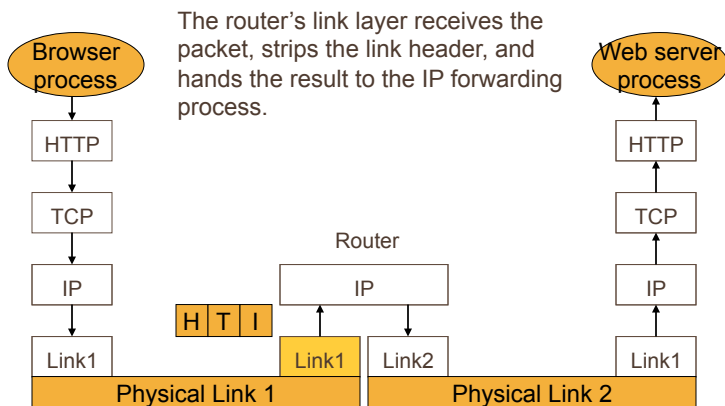
88

## Protocol layering example



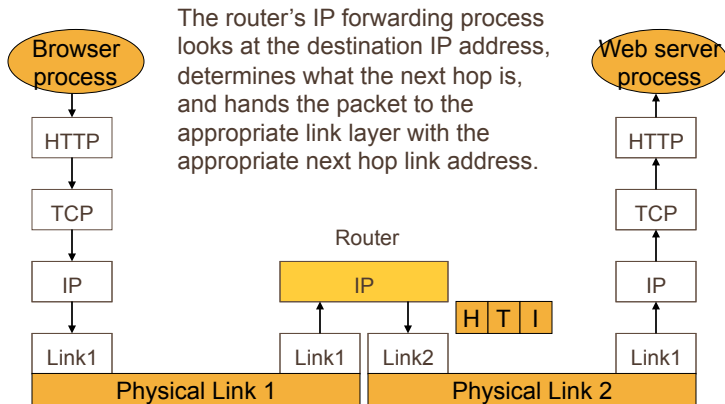
89

## Protocol layering example



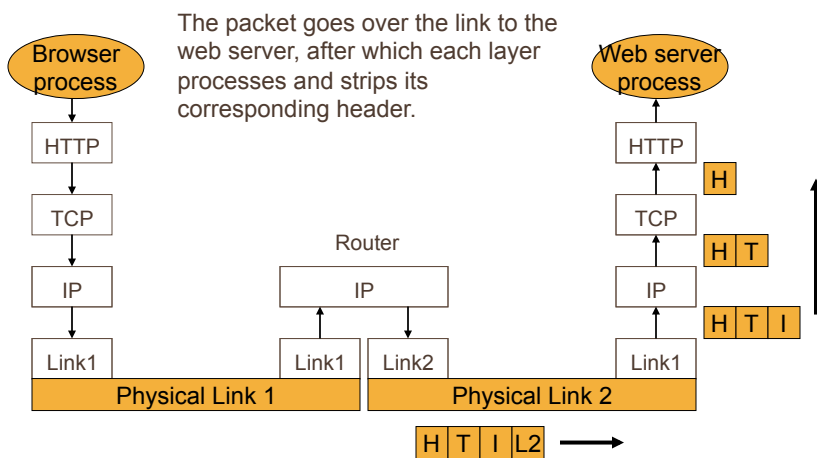
90

## Protocol layering example



91

## Protocol layering example



92

## TCP/IP PROTOCOLS

### Basic elements of any protocol header

- **Demuxing** field
  - Indicates which is the next higher layer (or process, or context, etc.)
- Length field or header delimiter
  - For the header, optionally for the whole packet
- Header format may be text (HTTP, SMTP (email)) or binary (IP, TCP, Ethernet)

## Demuxing fields

- **Ethernet:** Protocol Number
  - Indicates IPv4, IPv6, (old: Appletalk, SNA, Decnet, etc.)
- **IP:** Protocol Number
  - Indicates TCP, UDP, SCTP
- **TCP and UDP:** Port Number
  - Well known ports indicate FTP, SMTP, HTTP, SIP, many others
  - Dynamically negotiated ports indicate specific processes (for these and other protocols)
- **HTTP:** Host field
  - Indicates “virtual web server” within a physical web server
  - (Well, more like an identifier than a demuxing field)

95

## IP (Internet Protocol)

- Three services:
  - **Unicast:** transmits a packet to a specific host
  - **Multicast:** transmits a packet to a group of hosts
  - **Anycast:** transmits a packet to one of a group of hosts (typically nearest)
- Destination and source identified by the IP address (32 bits for IPv4, 128 bits for IPv6)
- All services are unreliable
  - Packet may be dropped, duplicated, and received in a different order

96



## IP address

- Both source and destination address may be modified in transit
  - By NAT boxes
  - But even so, can reply to the source IP address
  - Unless source address is spoofed
- IP (unicast) address is hierarchical, but host can treat it as a flat identifier (given net mask)
  - Can't tell how close or far a host is by looking at its IP address

97

## IP(v4) address format

- In binary, a 32-bit integer
- In text, this: “155.246.89.22”
  - Each decimal digit represents 8 bits (0 – 255)
- a.b.c.d/n
  - Last n bits identify machine
  - First 32-n bits identify network on the Internet
  - Ex: 155.246.89.0/8 for most Stevens hosts

98

## IP(v4) address format

- “Private” addresses are not globally unique:
  - Used behind NAT boxes
  - 10.0.0.0/8, 172.16.0.0/12, 192.168.0.0/16
- Multicast addresses start with 1110 as the first 4 bits (Class D address)
  - 224.0.0.0/4
- Unicast and anycast addresses come from the same space

99

## UDP (User Datagram Protocol)

- Runs above IP
- Same unreliable service as IP
  - Packets can get lost anywhere:
    - Outgoing buffer at source
    - Router or link
    - Incoming buffer at destination
- But adds port numbers
  - Mailboxes
  - Used to identify “application layer” processes
- Also a checksum, optional

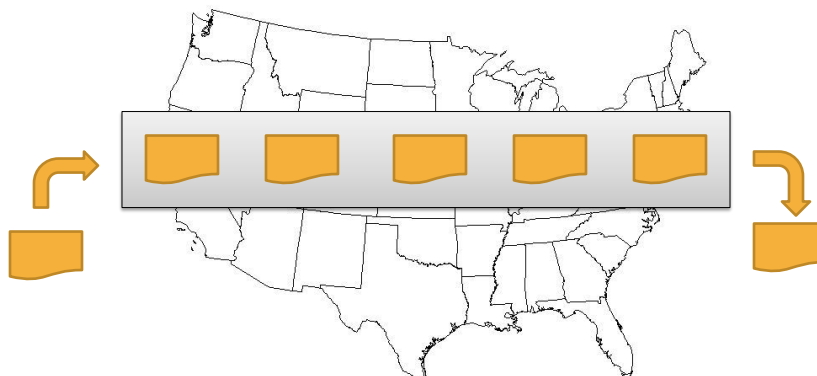
100

## TCP (Transmission Control Protocol)

- Runs above IP
  - Port number and checksum like UDP
- Service is in-order byte stream
  - Bytes transparently packaged in packets
- Flow control and congestion control
- Connection setup and teardown phases
- Can be considerable delay between bytes in at source and bytes out at destination
  - Because of timeouts and retransmissions
- Works only with unicast (not multicast or anycast)


101

## TCP Pipeline



102

## TCP Windowing



103

## UDP vs. TCP

- UDP is more real-time
  - Packet is sent or dropped, but is not delayed
- UDP has more of a “message” flavor
  - One packet = one message
  - But must add reliability mechanisms over it
- TCP good for transferring a file
  - Frustrating for messaging
  - Interrupts to application don't conform to message boundaries
  - No “Application Layer Framing”
- TCP is vulnerable to DoS (Denial of Service) attacks
  - SYN flood

104

## SCTP (Stream Control Transmission Protocol)

- IETF standard
- Overcomes many limitations of TCP
  - Motivation is SS7 signaling over IP
- Message oriented---supports message framing
- Multiple streams for a given session
  - Interruption in one stream does not effect others
- Cookie mechanism for DoS attacks

105

## Summary

- TCP, UDP, IP provide a nice set of basic tools
  - Key is to understand concept of protocol layering
- But problems/limitations exist
  - IP has been compromised by NAT, can't be used as a stable identifier
  - Firewalls can block communications
  - TCP has vulnerabilities
  - Network performance highly variable

106